

Theoretical Questions

Nitsan BenHanoach 208585927

I WAS ALLOWED TO SUBMIT THIS WORK LATE

Lior approved this because she is very nice <3

Part I - A

Determine whether the parentheses in a given mathematical equation are balanced.

Pseudo Code (C style)

```
char parenthesis_sign(char ch) {
    switch (ch) {
        case '(': return +1;
        case ')': return -1;
        default : return 0;
    }
}

bool are_parentheses_balanced(char expr[], int N) {
    int signs[N], balance[N];

    for (int i = 0; i < N; i++) do in parallel {
        signs[i] = parenthesis_sign(expr[i]);
    }

    balance = parallel_prefix_sum(signs);

    bool never_negative = true;

    for (int i = 0; i < N; i++) do in parallel {
        if (balance[i] < 0) {
            never_negative = false;
        }
    }

    return never_negative && balance[N - 1] == 0;
}
```

Explanation

The classic sequential algorithm scans the input left to right, incrementing a balance on every (, and decrementing it on every).

It's already pretty much a prefix-sum, if we treat (as +1 and) as -1; We just need to check that at no point the balance is negative, and that the final balance is zero.

Here we use parallel-prefix-sum as a black box, then verify (in parallel) that at no point the balance is negative.

Note that different threads assigning `never_negative = false` in parallel is *not* a race condition (there are no reads).

Correctness If the expression **is not balanced**, then **either** * there's a) with no a preceding (, or * there's a (with no succeeding).

In the first case, the balance at that) will be negative, and the algorithm will return false; In the second case, the *final* balance will be positive, and the algorithm will return false.

If the expression **is balanced**, then **both** * every (will increase the balance by 1, but it'll come with exactly one succeeding) that will decrease it by 1; so the final balance will be zero. * no) will come before its matching (, so at no point will the balance go negative.

Hence, the algorithm will return true.

Complexity In class we saw that `parallel_prefix_sum` can be done in $O(\log n)$ time.

My additions can be done in $O(1)$ (given N cores).

Hence `are_parentheses_balanced` can be run in $O(\log n)$ time.

(By the way, for the same reasons, it's also $O(n)$ work).

Part I - B

Find the number of continuous sub arrays with a given sum.

In the forum, Saar said we can assume all values are positive.

Pseudo Code (C style)

```
bool binary_search(int arr[], int N, int x) {
    // returns whether arr contains x
    // using classic sequential binary search
}
```

```

int continuous_sub_arr_sum(int arr[], int N, int target) {
    int arr_cumsum[N], results[N], results_cumsum[N];

    results[0] = (arr[0] == target) ? 1 : 0;

    arr_cumsum = parallel_prefix_sum(arr);

    for (int i = 1; i < N; i++) do in parallel {
        int complement = arr[i] - target;
        bool found = binary_search(arr, i, complement);
        results[i] = found ? 1 : 0;
    }

    results_cumsum = parallel_prefix_sum(results);

    return results_cumsum[N - 1];
}

```

Explanation

As an example, consider $*arr = [3, 2, 4, 5, 1, 9, 8, 1, 2, 6]$ $*target = 11$

After prefix-sum we get $*arr_cumsum = [3, 5, 9, 14, 15, 24, 32, 33, 35, 41]$

Notice $14 - 3 = 11$, and $35 - 24 = 11$.

We'll say that y is x 's *complement*, if $x - y == target$.

Each element in arr_cumsum can determine whether it has a preceding complement using binary search. Note that since $\forall i, arr[i] > 0$, arr_cumsum is strictly monotonous.

The first element is a special case (it has no preceding values), so we add a check for it.

But we still need to tackle **parallel mutual counting!**

- Letting threads increment a shared counter is a race condition.
- Protecting the counter with a mutex gives $O(n)$ runtime.

Solution: We create an array of results. Each thread sets their result to either 0 or 1. Afterwards, we sum the results (in parallel) in $O(\log n)$ time.

Complexity Assuming N threads, $parallel_prefix_sum(arr)$ takes $O(\log n)$ time.

The binary search takes another $O(\log n)$ for each thread in parallel.

The last $parallel_prefix_sum(results)$ again adds $O(\log n)$ time.

Total: $O(\log n)$ time.

(Work is $O(n \log n)$, because N threads perform logarithmic binary search).


Part II

Find the k -th ancestor of each node in an up-pointing tree.

Input: k , up-pointing tree (encoded as an array)

$k = 2$

indices: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9
array = [0, 0, 0, 1, 1, 2, 3, 5, 5, 5]



Output: k -th ancestor of each node, or -1 if no such

[-1, -1, -1, 0, 0, 0, 1, 2, 2, 2]

Pseudo Code (C style)

```
void parallel_copy_into(int dst[], int src[], int N) {
    for (int i = 0; i < N; i++) do in parallel {
        dst[i] = src[i];
    }
}

void deterministic_pointer_jumping(int arr[], int map[], int N) {
    // a deterministic version of pointer jumping,
    // working on a copy of arr, to avoid race conditions.

    int temp[N];

    for (int i = 0; i < N; i++) do in parallel {
        if (arr[i] > -1) {
            temp[i] = map[arr[i]];
        }
    }
}
```

```

    }

    parallel_copy_into(arr, temp, N);
}

void each_node_kth_ancestor(int arr[], int results[], int N) {

    //////////////////////////////////////
    // step 1: find root //
    //////////////////////////////////////

    int root;

    // copy arr into results
    parallel_copy_into(results, arr, N);

    // one round of pointer jumping is enough
    deterministic_pointer_jumping(results, results, N);

    // root will be the only node pointing to itself
    for (int i = 0; i < N; i++) do in parallel {
        if (results[i] == i) {
            root = i;
        }
    }

    //////////////////////////////////////
    // step 2: make root point nowhere //
    //////////////////////////////////////

    arr[root] = -1;

    //////////////////////////////////////
    // step 3: money time //
    //////////////////////////////////////

    // copy arr into results
    parallel_copy_into(results, arr, N);

    while (k) {
        if (k & 0x1) {
            // pointer jumping in results based on arr
            deterministic_pointer_jumping(results, arr, N);
        }
        // pointer jumping in arr
        deterministic_pointer_jumping(arr, arr, N);
    }
}

```

```

        k = k >> 1;
    }
}

```

Explanation

My approach takes inspiration from the famous Exponentiation by Squaring algorithm.

```

double fast_pow(double x, unsigned int n) {
    double result = 1.0;
    while (n) {
        if (n & 0x1) {
            result *= x;
        }
        x *= x;
        n >>= 1;
    }
    return result;
}

```

We perform regular pointer jumping in **arr on every iteration**. * At iteration *i*, **arr** has links of distance $2^{**} i$.

We perform *guided* pointer jumping in **results on some iterations**. * We only jump if the *i*'th bit of *k* is set. * We advance each result by a distance of $2^{**} i$ (using **arr**).

Example Let's say *k* = 43; its binary representation is 0b101011, accounting to the fact that $43 = 32 + 0 + 8 + 0 + 2 + 1$.

i	arr	results
0	jumps of 1	jump += 1
1	jumps of 2	jump += 2
2	jumps of 4	
3	jumps of 8	jump += 8
4	jumps of 16	
5	jumps of 32	jump += 32

Correctness Normal pointer jumping (as in class) corresponds with finding ancestors at distance of **at least** *k*.

However, in our case we need to find ancestors at distance of **exactly** *k*. Also, nodes without a *k*'th ancestor should point to -1.

In this algorithm, we start by we making the root point to -1. We also treat -1 as a black whole (once a node gets to -1, it'll stay there forever).

The trick of selectively using powers of 2, with -1 being a black whole, guarantees we'll indeed find **exact** k 'th ancestors.

Complexity Assuming N cores.

`parallel_copy_into` takes $O(1)$ time.

`deterministic_pointer_jumping` takes $O(1)$ time.

Step 1 does Pointer-Jumping once, so it's $O(1)$ time.

Step 2 is $O(1)$ time.

Step 3 does Pointer-Jumping for every bit of k , so it's $O(\log k)$.

Total: $O(\log k)$ time.